

BUS 41201 Homework 4 Assignment

Group 24: Shihan Ban, Yi Cao, Shri Lekkala, Ningxin Zhang

16 April 2024

Setup

```
## microfinance network
## data from BANERJEE, CHANDRASEKHAR, DUFLO, JACKSON 2012

## data on 8622 households
hh <- read.csv("microfi_households.csv", row.names="hh")
hh$village <- factor(hh$village)

## We'll kick off with a bunch of network stuff.
## This will be covered in more detail in lecture 6.
## get igraph off of CRAN if you don't have it
## install.packages("igraph")
## this is a tool for network analysis
## (see http://igraph.sourceforge.net/)
library(igraph)
```

```
##
## Attaching package: 'igraph'

## The following objects are masked from 'package:stats':
##
##      decompose, spectrum

## The following object is masked from 'package:base':
##
##      union
```

```
edges <- read.table("microfi_edges.txt", colClasses="character")
## edges holds connections between the household ids
hhnet <- graph.edgelist(as.matrix(edges))
hhnet <- as.undirected(hhnet) # two-way connections.

## igraph is all about plotting.
V(hhnet) ## our 8000+ household vertices
```

```
## + 8182/8182 vertices, named, from 128e388:
##      [1] 1002 1001 1020 1042 1053 1163 1003 1004 1026 1029 1076 1159
##      [13] 1106 1031 1048 1081 1006 1005 1008 1016 1021 1024 1089 1103
##      [25] 1007 1019 1155 1015 1040 1044 1045 1078 1088 1110 1115 1140
##      [37] 1145 1009 1018 1060 1064 1073 1153 1067 1099 1010 1162 1012
##      [49] 1143 1013 1023 1028 1034 1065 1117 1139 1154 1157 1173 1014
##      [61] 1068 1071 1148 1017 1036 1062 1112 1118 1120 1129 1134 1165
##      [73] 1183 1126 1122 1049 1058 1093 1108 1114 1119 1022 1043 1079
```


```
##      [85] 1033  1102  1104  1105  1152  1169  1171  1025  1027  1147  1032  1035
##      [97] 1037  1039  1041  1113  1174  1069  1116  1132  1178  1146  1080  1086
##     [109] 1101  1172  1059  1141  1142  1038  1094  1052  1092  1082  1095  1158
## + ... omitted several vertices
```

```
## Each vertex (node) has some attributes, and we can add more.
V(hhnet)$village <- as.character(hh[V(hhnet),'village'])
## we'll color them by village membership
vilcol <- rainbow(nlevels(hh$village))
names(vilcol) <- levels(hh$village)
V(hhnet)$color = vilcol[V(hhnet)$village]
## drop HH labels from plot
V(hhnet)$label=NA


# graph plots try to force distances proportional to connectivity
# imagine nodes connected by elastic bands that you are pulling apart
# The graphs can take a very long time, but I've found
# edge.curved=FALSE speeds things up a lot. Not sure why.

## we'll use induced.subgraph and plot a couple villages
village1 <- induced.subgraph(hhnet, v=which(V(hhnet)$village=="1"))
village33 <- induced.subgraph(hhnet, v=which(V(hhnet)$village=="33"))

# vertex.size=3 is small. default is 15
plot(village1, vertex.size=3, edge.curved=FALSE)
```



```
plot(village33, vertex.size=3, edge.curved=FALSE)
```



```
library(gamlr)
```

```
## Loading required package: Matrix
```

```
## match id's; I call these 'zebras' because they are like crosswalks
zebra <- match(rownames(hh), V(hhnet)$name)

## calculate the `degree' of each hh:
## number of commerce/friend/family connections
degree <- degree(hhnet)[zebra]
names(degree) <- rownames(hh)
degree[is.na(degree)] <- 0 # unconnected houses, not in our graph


## if you run a full glm, it takes forever and is an overfit mess
# > summary(full <- glm(loan ~ degree + .^2, data=hh, family="binomial"))
# Warning messages:
# 1: glm.fit: algorithm did not converge
# 2: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

Question 1

I'd transform degree to create our treatment variable d. What would you do and why?

We can first plot a histogram of the degree variable to get an idea of its structure:

```
hist(degree)
```



From the graph, it might be appropriate to perform a logarithmic transformation for the following reasons:

- It appears that the degree frequency is highly skewed to the right as there are many nodes with few connections (degree < 20), but few nodes with many connections (degree > 40). So by taking a log transformation, we can normalize the distribution, making it more symmetric and more suitable to statistical analyses.
- The histogram appears to follow an exponential / multiplicative relationship. So transforming the data logarithmically can make the relationship more linear, which is easier to model and interpret in regression models.
- We can reduce the range of variability in degree values, effectively performing a dimensionality reduction. This is useful to prevent the model being overly effected by outliers, i.e. households with a very high number of connections.

```
# Transform degree and add it to the hh dataset
hh$d = log1p(degree)
head(hh)
```

```
##      loan village religion  roof rooms beds electricity ownership leader
## 1001    0       1    hindu  tile     3   4           0     OWNED      0
## 1002    0       1    hindu  tile     1   1           1     OWNED      1
## 1003    0       1    hindu  rcc     3   4           1     OWNED      1
## 1004    0       1    hindu  tile     2   6           1     OWNED      0
```

##	1005	0	1	hindu tile	3	4	1	OWNED	0
##	1006	0	1	hindu stone	2	1	1	OWNED	0
##				d					
##	1001	1.791759							
##	1002	2.079442							
##	1003	1.098612							
##	1004	1.609438							
##	1005	2.197225							
##	1006	2.302585							

Question 2

Build a model to predict d from x , our controls.


Comment on how tight the fit is, and what that implies for estimation of a treatment effect.

```
model_d <- lm(d ~ village + religion + roof + rooms + beds + electricity + ownership + leader, data = hh)
# Display model summary
summary(model_d)
```


```
##
## Call:
## lm(formula = d ~ village + religion + roof + rooms + beds + electricity +
##      ownership + leader, data = hh)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.6311 -0.5365  0.1161  0.6385  2.2392
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.473572   0.325760   4.523 6.16e-06 ***
## village2       -0.242852   0.086019  -2.823 0.004765 **
## village3       -0.058747   0.079540  -0.739 0.460182
## village4       -0.054250   0.083099  -0.653 0.513881
## village12      -0.059632   0.090571  -0.658 0.510299
## village19       0.060352   0.084849   0.711 0.476925
## village20      -0.321265   0.090808  -3.538 0.000406 ***
## village21       0.113996   0.085170   1.338 0.180784
## village23      -0.059661   0.081211  -0.735 0.462579
## village24       0.114800   0.090560   1.268 0.204952
## village25       0.008783   0.081591   0.108 0.914277
## village28      -0.070921   0.078517  -0.903 0.366417
## village29      -0.047275   0.079156  -0.597 0.550365
## village31      -0.004052   0.093999  -0.043 0.965617
## village32       0.056103   0.082910   0.677 0.498630
## village33      -0.140952   0.085098  -1.656 0.097689 .
## village36       0.051463   0.078448   0.656 0.511830
## village39      -0.094507   0.079081  -1.195 0.232094
## village42      -0.188036   0.087883  -2.140 0.032414 *
## village43      -0.091643   0.087457  -1.048 0.294728
## village45      -0.104098   0.084104  -1.238 0.215848
## village46       0.030666   0.083264   0.368 0.712662
## village47       0.046783   0.094194   0.497 0.619436
## village50      -0.008535   0.082470  -0.103 0.917572
## village51       0.191341   0.082183   2.328 0.019923 *
## village52       0.212353   0.077862   2.727 0.006398 **
## village55      -0.217925   0.082782  -2.633 0.008491 **
## village57      -0.009889   0.085943  -0.115 0.908400
## village59      -0.067312   0.078466  -0.858 0.391001
## village62       0.059493   0.086566   0.687 0.491939
## village65       0.178011   0.079319   2.244 0.024842 *
```

```
## village67      0.200421  0.086384  2.320 0.020359 *
## village68      0.109751  0.091862  1.195 0.232222
## village70      0.307165  0.085385  3.597 0.000323 ***
## village71      0.122435  0.079403  1.542 0.123122
## village72      0.181744  0.083557  2.175 0.029652 *
## village73      0.277864  0.088691  3.133 0.001736 **
## village75      0.209743  0.089709  2.338 0.019408 *
## religionhindu   0.125830  0.312849  0.402 0.687541
## religionislam  -0.029901  0.315721 -0.095 0.924551
## roofrcc        -0.025928  0.051691 -0.502 0.615972
## roofsheet      -0.062450  0.047527 -1.314 0.188884
## roofstone      -0.007146  0.046165 -0.155 0.876987
## roofthatch     -0.039734  0.075249 -0.528 0.597490
## rooftile       -0.008428  0.047101 -0.179 0.857998
## rooms          0.072171  0.008184  8.819 < 2e-16 ***
## beds           0.018469  0.007467  2.473 0.013403 *
## electricity     0.265259  0.036576  7.252 4.45e-13 ***
## ownershipLEASED -0.509224  0.220214 -2.312 0.020779 *
## ownershipOWNED  -0.019337  0.052808 -0.366 0.714246
## ownershipRENTED -0.256670  0.065819 -3.900 9.71e-05 ***
## ownershipSHARE_OWNED -0.217944  0.108794 -2.003 0.045178 *
## leader         0.310300  0.027918 11.115 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8246 on 8569 degrees of freedom
## Multiple R-squared:  0.08223,    Adjusted R-squared:  0.07666
## F-statistic: 14.76 on 52 and 8569 DF,  p-value: < 2.2e-16
```

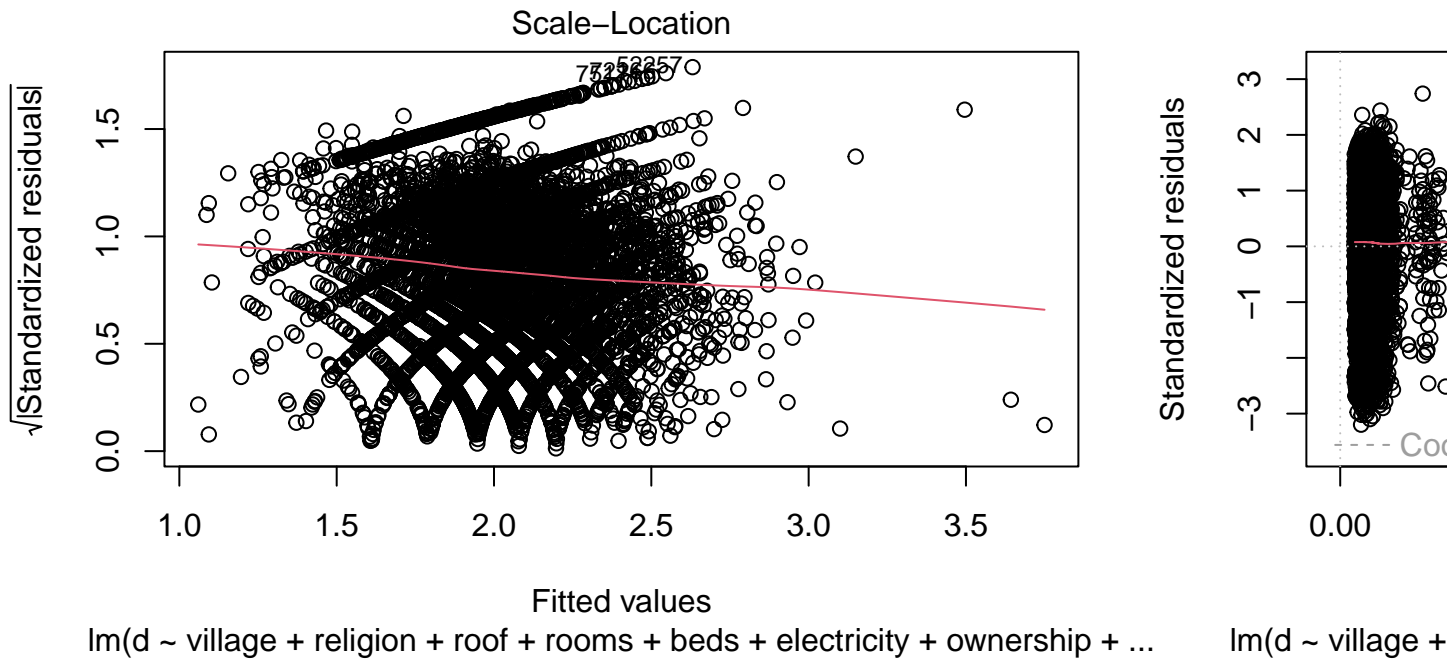
```
plot(model_d)
```



lm(d ~ village + religion + roof + rooms + beds + electricity + ownership + ...



lm(d ~ village +



In this model, some variables, like rooms, electricity, beds, and leader, are statistically significant, implying they have a notable impact on the response variable.

The model we've run seems reasonably well-specified based on the first plot but does show potential outliers. The presence of influential points, as suggested by the fourth plot, means that the estimated treatment effect may be driven by a small number of cases. Thus, while there are indicators of a decent model fit, the presence of outliers and influential points suggests that the treatment effect estimation may be sensitive to particular data points. When interpreting treatment effects, it's crucial to consider these factors, as they can affect the validity and generalizability of the findings.


```

# control variables
x = model.matrix(d ~ village + religion + roof + rooms + beds + electricity + ownership + leader - 1, d

# dependent variable
y = hh$loan


# treatment variable
d = hh$d

# Estimate d_hat with lasso regression of d on x.
treat = gamlr(x, d, lambda.min.ratio=1e-4)

# Isolate dhat (the part of treatment that we can predict with x's)
d_hat = predict(treat, x, type="response")

# Plot d_hat against d
plot(d_hat, d)

```



To assess the model fit, we can compute the in sample R2

```

# In-sample R2
cor(drop(d_hat), d)^2

```

```
## [1] 0.08187873
```

So the in-sample R2 value suggests that $\approx 8.19\%$ of the variance in d is explained by the control variables. Thus the model does not have a tight fit, and this implies that there may be other confounding variables not included in the model that account for d .

Thus the predictive power of our model is limited due to the large percentage of unexplained variance in d . And further analyses may lead to less accurate and biased estimates, so our confidence in a estimating treatment effect would be low.

Question 3

Use predictions from [2] in an estimator for effect of d on loan.

```
# Second Stage Lasso  
  
# Do a lasso of y on [d, d_hat, x], with d_hat unpenalized  
causal = gamlr(cbind(d, d_hat, x), y, free=2)
```

```
## 'as(<dgeMatrix>, "dgCMatrix")' is deprecated.  
## Use 'as(., "CsparseMatrix")' instead.  
## See help("Deprecated") and help("Matrix-deprecated").
```

```
# Second  
print(coef(causal)["d",])
```

```
## [1] 0.0187176
```

Using the two-stage lasso process, we find the best predictor for y from d and x after the influence of d_hat is removed.

We observe that the coefficient of the log transformed degree variable is 0.0187176, which suggests that there is a small positive relationship between the degree of connectivity and the likelihood of adopting a loan.

```
exp(coef(causal)["d",])
```

```
## [1] 1.018894
```

By taking the exponential of the coefficient, we compute the odds ratio between degree and loan. That is, a one unit increase in the log transformed degree of connection corresponds to an $\approx 1.89\%$ increase in the probability of a household taking a loan.

Question 4

Compare the results from [3] to those from a straight (naive) lasso for loan on \mathbf{d} and \mathbf{x} .

```
# Compute a naive lasso for loan
# We use binomial here since we know y (loan) is in [0,1]
naive = gamlr(cbind(d, x), y, family="binomial")

# Compare naive and 2-stage lasso
cat("The coefficient for d from the naive lasso is:", coef(naive)["d",], "\n")
```

```
## The coefficient for d from the naive lasso is: 0.1562872
```

```
cat("The coefficient for d from the causal lasso is:", coef(causal)["d",])
```

```
## The coefficient for d from the causal lasso is: 0.0187176
```

```
exp(coef(naive)["d",])
```

```
## [1] 1.169162
```

Explain why they are similar or different.

Firstly we observe that the coefficient for \mathbf{d} from the naive model is approximately an order of 10 greater than the one from the causal model. That is, a one unit increase in log transformed degree would suggest there would be a $\approx 16.9\%$ increase in the probability of a household taking a loan (compared to $\approx 1.89\%$ from Q3).

They are different as the naive model does not separate the treatment and the control variables, but rather uses them all as independent variables in the regression. This may result in the coefficient for \mathbf{d} having contributions from confounding variables which are not accounted for, and thus indicates a much more significant effect than the causal model.

In comparison, the causal model involved a 2-stage LASSO process by incorporating $\mathbf{d_hat}$, the predicted values of \mathbf{d} based on \mathbf{x} . This means that this model controlled for the portion of \mathbf{d} that could be predicted from the control variables, aiming to isolate the more variation in \mathbf{d} . (However there could still be effects from confounding variables that were not in the data set).

So the two-stage Lasso model provides a more conservative but likely more accurate estimate by explicitly modeling and removing the predictable part of \mathbf{d} based on the observed covariates \mathbf{x} .

Question 5

Bootstrap your estimator from [3] and describe the uncertainty.

[+]

Can you think of how you'd design an experiment to estimate the treatment effect of network degree?